

从零开始的深度学习课程

II 工具解析

谢晋璟


Commerce Data Science Team

CDL数据科学社区（原BDA社区）
数据科学工作室
技术茶会

2016-07-22

窥视深渊者,必为深渊所窥视

And when you look long into an abyss, the abyss also looks into you. - Friedrich Nietzsche



目录

- 1 显卡计算简介
- 2 Caffe框架解析
 - Blob—数据对象
 - Layer—变换层
 - Net—整体网络
 - Solver—求解器
- 3 TensorFlow—兵器谱之首位
 - TensorFlow为什么叫TensorFlow?
 - TensorFlow进行简单计算例子
 - Perceptron的例子

一个比喻：GPU vs 集群

GPU



集群



	Throughput	Latency
GPU	High	High
CPU	Low	Low

Table: GPU和CPU对比

GPU计算能力- 例子

桌面- GeForce GTX 1080

- Double Flops: NA
- Single Flops: 8.2T
- CUDA core: 2560
- Compute Capability: 6.1
- GPU Architecture: Pascal
- GPU Memory: 8GB
- Price: about 5000¥

服务器/工作站- Tesla K40

- Double Flops: 1.43T
- Single Flops: 4.29T
- CUDA core: 2880
- Compute Capability: 3.5
- GPU Architecture: Kepler
- GPU Memory: 12GB
- Price: around 20000¥

对比CPU的flops，比如i7-6700K(超频到4.6GHz)最多只达到211GFlops。一些简单Layer（比如Affine）不需要复杂的指令集，分支预测，乱序执行等现代CPU的高级功能，但它是强计算密集型的，正是发挥GPU高能耗比的地方。

Show me the code – “Hello, World”

这是CUDA下的“Hello, World”–向量加法，计算

$$\text{out} \leftarrow \text{in1} + \text{in2}$$

下面是设备（CUDA只支持GPU，OPENCL理论上还支持其他设备FPGA等）端代码

Device Side – Kernel Code

```
__global__ void vecAdd(float *in1, float *in2, float *out, int len) {  
    int i = threadIdx.x + blockDim.x * blockIdx.x ;  
    if (i < len)    out[i] = in1[i] + in2[i] ;  
}
```

下一页是主机端代码

主机端代码

```
//apply for device memory
int size = sizeof(float) * inputLength;
cudaMalloc((void **) &deviceInput1, size);
cudaMalloc((void **) &deviceInput2, size);
cudaMalloc((void **) &deviceOutput, size);
//copy data to device
cudaMemcpy(deviceInput1, hostInput1, size, cudaMemcpyHostToDevice);
cudaMemcpy(deviceInput2, hostInput2, size, cudaMemcpyHostToDevice);
//define grid for computation
dim3 DimGrid((inputLength-1)/256 + 1, 1, 1);
dim3 DimBlock(256 , 1, 1);
//invoke the Kernel for computation
vecAdd<<<DimGrid, DimBlock>>>(deviceInput1, deviceInput2, deviceOutput,
    inputLength);
//need to synchronize between host and device
cudaDeviceSynchronize();
//After device finished its computation, copy the results back
cudaMemcpy(hostOutput, deviceOutput, size, cudaMemcpyDeviceToHost);
//release the device resources
cudaFree(deviceInput1); cudaFree(deviceInput2);
cudaFree(deviceOutput);
```

Bandwidth的诅咒

上面的例子里计算和读写内存的比例是1:1, 这个比值称为

CGMA(compute to global memory access ratio)

比如NVIDIA G80只有86.4GB/s带宽, 最多调入21.6G的单精度浮点数, 但CGMA=1时, 最多只能执行21.6GFlop的浮点计算, 远远小于设备的峰值吞吐量367GFlops。

如何逼近Throughput峰值?

解决办法

使用片上shared memory

其他影响Bandwidth的内容

合并访存


NVIDIA的库

编写GPU计算程序要考虑的细节非常繁多，需要对硬件和编译器有较深入的理解才能写出高效的代码，不然速度可能还不如CPU。好在已经有了很多现成的library可以使用，大部分情况我们只要调用一下，不必从头造轮子。比如NVIDIA自己就提供了这些

- **cuFFT** Fast Fourier Transforms Library
- **cuBLAS** Complete BLAS library
- **cuSolver** Collection of dense and sparse direct solvers/cusolver
- **cuSPARSE** Sparse Matrix library
- **cuRAND** Random Number Generator
- **NPP** Thousands of Performance Primitives for Image & Video Processing
- **Thrust** Templated Parallel Algorithms & Data Structures
- **CUDA Math Library** high performance math routines
- **cuDNN** library for deep neural networks

现有DL工具对GPU的需求

回到我们的主题，所有的DL工具都默认使用GPU。
最常用的DL工具就是Caffe, Torch7, Theano和TensorFlow。
下面借了大MS的一张广告图：





**Decaf
Coffee
is Caffeine
Free**

前言

Caffe是最早流行的DL工具，它基本上只是下面两个东西

- C++ 写的library
- 命令行下的训练、预测工具

后来它也提供了python的接口，可以在python中比较方便的存取它的数据对象，调用向前向后命令等。

Caffe的代码量不大，写的比较清晰，和我们在上一讲中的介绍是一一对应的，所以它是个理解DL框架的完美例子。下面一一解释它的核心对象。

目录

- 1 显卡计算简介
- 2 Caffe框架解析
 - Blob—数据对象
 - Layer—变换层
 - Net—整体网络
 - Solver—求解器
- 3 TensorFlow—兵器谱之首位
 - TensorFlow为什么叫TensorFlow?
 - TensorFlow进行简单计算例子
 - Perceptron的例子

Blob是存储张量的数据结构

Blob的域

```
protected:
    shared_ptr<SyncedMemory> data_;           // X
    shared_ptr<SyncedMemory> diff_;           //  $\delta X = \frac{\partial l}{\partial X}$ 
    shared_ptr<SyncedMemory> shape_data_;     // same as shape_
    vector<int> shape_;                        // dimension vector
    int count_;                                // number of elements
    int capacity_;
```

基础数据对象Blob

用来存储所有的（张量）数据

- 输入数据和神经元的激活: $X_{u,v}^{c,j}$
- Score: S
- 反馈信号: $\delta X_{u,v}^{c,j}, \delta S$
- 待估参数: $W_{u,v}^l, K(l)(u, v)_c'$

Blob是绑定了data和loss对此data偏导的数据结构

它的update方法就是用自己的d X修正自己的X.

Blob的update方法

$$W \leftarrow W + (-1) \cdot dW$$

```
template <typename Dtype>
void Blob<Dtype>::Update() {
    // We will perform update based on where the data is located.
    switch (data_>head()) {
        case SyncedMemory::HEAD_AT_CPU:
            // perform computation on CPU
            caffe_axpy<Dtype>(count_, Dtype(-1),
                static_cast<const Dtype*>(diff_>cpu_data()),
                static_cast<Dtype*>(data_>mutable_cpu_data()));
            break;
        case SyncedMemory::HEAD_AT_GPU:
        case SyncedMemory::SYNCED:
#ifdef CPU_ONLY
            // perform computation on GPU
            caffe_gpu_axpy<Dtype>(count_, Dtype(-1),
                static_cast<const Dtype*>(diff_>gpu_data()),
                static_cast<Dtype*>(data_>mutable_gpu_data()));
            // . . .
    }
```

目录

- 1 显卡计算简介
- 2 Caffe框架解析
 - Blob—数据对象
 - Layer—变换层
 - Net—整体网络
 - Solver—求解器
- 3 TensorFlow—兵器谱之首位
 - TensorFlow为什么叫TensorFlow?
 - TensorFlow进行简单计算例子
 - Perceptron的例子

Layer就是我们的多元向量映射 $f(X; W)$

Layer类的域

```
protected:
  /** The protobuf that stores the layer parameters */
  LayerParameter layer_param_;
  /** The phase: TRAIN or TEST */
  Phase phase_;
  /** The vector that stores the learnable parameters as a set of blobs. */
  vector<shared_ptr<Blob<Dtype> > > blobs_;
  /** Vector indicating whether to compute the diff of each param blob. */
  vector<bool> param_propagate_down_;

  /** The vector that indicates whether each top blob has a non-zero weight in
   * the objective function. */
  vector<Dtype> loss_;
```

其中，blobs_ 用于存储 W ，而LayerParameter 才是记录输入输出，类型的地方。
它是由protobuf定义的，定义文件的一部分如下：

LayerParameter

```
message LayerParameter {
  optional string name = 1; // the layer name
  optional string type = 2; // the layer type
  repeated string bottom = 3; // the name of each bottom blob
  repeated string top = 4; // the name of each top blob

  // The train / test phase for computation.
  optional Phase phase = 10;

  // The amount of weight to assign each top blob in the objective.
  // Each layer assigns a default value, usually of either 0 or 1,
  // to each top blob.
  repeated float loss_weight = 5;

  // Specifies training parameters (multipliers on global learning constants,
  // and the name and other settings used for weight sharing).
  repeated ParamSpec param = 6;

  // The blobs containing the numeric parameters of the layer.
  repeated BlobProto blobs = 7;

  // Specifies whether to backpropagate to each bottom. If unspecified,
  // Caffe will automatically infer whether each input needs backpropagation
  // to compute parameter gradients. If set to true for some inputs,
  // backpropagation to those inputs is forced; if set false for some inputs,
  // backpropagation to those inputs is skipped.
  //
  // The size must be either 0 or equal to the number of bottoms.
  repeated bool propagate_down = 11;
```

LayerParameter (续)

```
// Rules controlling whether and when a layer is included in the network,
// based on the current NetState.  You may specify a non-zero number of rules
// to include OR exclude, but not both.  If no include or exclude rules are
// specified, the layer is always included.  If the current NetState meets
// ANY (i.e., one or more) of the specified rules, the layer is
// included/excluded.
repeated NetStateRule include = 8;
repeated NetStateRule exclude = 9;

// Parameters for data pre-processing.
optional TransformationParameter transform_param = 100;

// Parameters shared by loss layers.
optional LossParameter loss_param = 101;

// . . .
optional ConvolutionParameter convolution_param = 106;
optional InnerProductParameter inner_product_param = 117;
optional InputParameter input_param = 143;
optional ReLUParameter relu_param = 123;
optional SoftmaxParameter softmax_param = 125;
optional DataParameter data_param = 107;
// . . .
```

它包含了具体的Layer的parameter。根据type不同，里面出现的可能是convolution_param, relu_param等，下面举个ConvolutionParameter的例子：

ConvolutionParameter(例子)

```
message ConvolutionParameter {
  optional uint32 num_output = 1; // The number of outputs for the layer
  optional bool bias_term = 2 [default = true]; // whether to have bias terms

  // Pad, kernel size, and stride are all given as a single value for equal
  // dimensions in all spatial dimensions, or once per spatial dimension.
  repeated uint32 pad = 3; // The padding size; defaults to 0
  repeated uint32 kernel_size = 4; // The kernel size
  repeated uint32 stride = 6; // The stride; defaults to 1
  // Factor used to dilate the kernel, (implicitly) zero-filling the resulting
  // holes. (Kernel dilation is sometimes referred to by its use in the
  // algorithm trous from Holschneider et al. 1987.)
  repeated uint32 dilation = 18; // The dilation; defaults to 1

  // For 2D convolution only, the *_h and *_w versions may also be used to
  // specify both spatial dimensions.
  optional uint32 pad_h = 9 [default = 0]; // The padding height (2D only)
  optional uint32 pad_w = 10 [default = 0]; // The padding width (2D only)
  optional uint32 kernel_h = 11; // The kernel height (2D only)
  optional uint32 kernel_w = 12; // The kernel width (2D only)
  optional uint32 stride_h = 13; // The stride height (2D only)
  optional uint32 stride_w = 14; // The stride width (2D only)

  optional uint32 group = 5 [default = 1]; // The group size for group conv

  optional FillerParameter weight_filler = 7; // The filler for the weight
  optional FillerParameter bias_filler = 8; // The filler for the bias
```

ConvolutionParameter(续)

```
enum Engine {  
    DEFAULT = 0;  
    CAFFE = 1;  
    CUDNN = 2;  
}  
optional Engine engine = 15 [default = DEFAULT];  
  
// The axis to interpret as "channels" when performing convolution.  
// Preceding dimensions are treated as independent inputs;  
// succeeding dimensions are treated as "spatial".  
// With (N, C, H, W) inputs, and axis == 1 (the default), we perform  
// N independent 2D convolutions, sliding C-channel (or (C/g)-channels, for  
// groups g>1) filters across the spatial axes (H, W) of the input.  
// With (N, C, D, H, W) inputs, and axis == 1, we perform  
// N independent 3D convolutions, sliding (C/g)-channels  
// filters across the spatial axes (D, H, W) of the input.  
optional int32 axis = 16 [default = 1];  
  
// Whether to force use of the general ND convolution, even if a specific  
// implementation for blobs of the appropriate number of spatial dimensions  
// is available. (Currently, there is only a 2D-specific convolution  
// implementation; for input blobs with num_axes != 2, this option is  
// ignored and the ND implementation will be used.)  
optional bool force_nd_im2col = 17 [default = false];  
}
```

可以看到上一讲提到的概念padding, stride等都出现在里面。

DataParameter—负责数据的输入的也被当成Layer

```
message DataParameter {  
  // Specify the data source.  
  optional string source = 1;  
  // Specify the batch size.  
  optional uint32 batch_size = 4;  
  
  optional DB backend = 8 [default = LEVELDB];  
  
  optional string mean_file = 3;  
  
  // Force the encoded image to have 3 color channels  
  optional bool force_encoded_color = 9 [default = false];  
  
  // Prefetch queue (Number of batches to prefetch to host memory, increase if  
  // data access bandwidth varies).  
  optional uint32 prefetch = 10 [default = 4];  
}
```

DataLayer是非常重要的**Layer**，它虽然不负责计算，也没有向后方法，但是它的**Forward**方法负责读取一个**Batch**的数据。

Layer的最重要方法：向前和向后

下面是它的Interface:

```
/**
 * The Forward wrapper calls the relevant device wrapper function
 * (Forward_cpu or Forward_gpu) to compute the top blob values given the
 * bottom blobs. If the layer has any non-zero loss_weights, the wrapper
 * then computes and returns the loss.
 *
 * Your layer should implement Forward_cpu and (optionally) Forward_gpu.
 */
inline Dtype Forward(const vector<Blob<Dtype>*>& bottom,
                    const vector<Blob<Dtype>*>& top);

/**
 * @param propagate_down
 *   a vector with equal length to bottom, with each index indicating
 *   whether to propagate the error gradients down to the bottom blob at
 *   the corresponding index
 *
 * The Backward wrapper calls the relevant device wrapper function
 * (Backward_cpu or Backward_gpu) to compute the bottom blob diffs given the
 * top blob diffs.
 *
 * Your layer should implement Backward_cpu and (optionally) Backward_gpu.
 */
inline void Backward(const vector<Blob<Dtype>*>& top,
                    const vector<bool>& propagate_down,
                    const vector<Blob<Dtype>*>& bottom);
```

目录

- 1 显卡计算简介
- 2 Caffe框架解析
 - Blob—数据对象
 - Layer—变换层
 - Net—整体网络
 - Solver—求解器
- 3 TensorFlow—兵器谱之首位
 - TensorFlow为什么叫TensorFlow?
 - TensorFlow进行简单计算例子
 - Perceptron的例子

Net就是hypotheses

Net把各个计算用的Layer的输入输出和参数按照一定的方式组织在一起。它负责计算loss值和loss的对参数的梯度。下面详解它的类定义。

Net的域(1): 层和块

```
/// @brief The network name
string name_;
/// @brief The phase: TRAIN or TEST
Phase phase_;

/// @brief Individual layers in the net
vector<shared_ptr<Layer<Dtype> > > layers_;
vector<string> layer_names_;
map<string, int> layer_names_index_;
vector<bool> layer_need_backward_;

/// @brief the blobs storing intermediate results between the layer.
vector<shared_ptr<Blob<Dtype> > > blobs_;
vector<string> blob_names_;
map<string, int> blob_names_index_;
vector<bool> blob_need_backward_;
```

可以发现它是由Layer和描述激活度的Blob连接成的。

Net的域(2): 每层的输入输出, 整体的输入输出块

```
/// bottom_vecs stores the vectors containing the input for each layer.  
/// They don't actually host the blobs (blobs_ does), so we simply store  
/// pointers.  
vector<vector<Blob<Dtype>*> > bottom_vecs_;  
vector<vector<int> > bottom_id_vecs_;  
vector<vector<bool> > bottom_need_backward_;  
  
/// top_vecs stores the vectors containing the output for each layer  
vector<vector<Blob<Dtype>*> > top_vecs_;  
vector<vector<int> > top_id_vecs_;  
  
/// Vector of weight in the loss (or objective) function of each net blob,  
/// indexed by blob_id.  
vector<Dtype> blob_loss_weights_;  
  
/// blob indices for the input and the output of the net  
vector<int> net_input_blob_indices_;  
vector<int> net_output_blob_indices_;  
vector<Blob<Dtype>*> net_input_blobs_;  
vector<Blob<Dtype>*> net_output_blobs_;
```

它们都是一些book keeping的指针, 记录哪些Blob是哪些Layer的输入输出, 以及整体的输入是哪个Blob, 输出是哪个Blob。

Net的域(3): 参数全体—固定的和可学习的

```
vector<vector<int> > param_id_vecs_;
vector<int> param_owners_;
vector<string> param_display_names_;
vector<pair<int, int> > param_layer_indices_;
map<string, int> param_names_index_;

/// The parameters in the network.
vector<shared_ptr<Blob<Dtype> > > params_;
vector<Blob<Dtype>*> learnable_params_;
/**
 * The mapping from params_ -> learnable_params_: we have
 * learnable_param_ids_.size() == params_.size(),
 * and learnable_params_[learnable_param_ids_[i]] == params_[i].get()
 * if and only if params_[i] is an "owner"; otherwise, params_[i] is a sharer
 * and learnable_params_[learnable_param_ids_[i]] gives its owner.
 */
vector<int> learnable_param_ids_;
```

Net的域(4): 调节网络在训练中的行为的参数, learning rate等

```
/// the learning rate multipliers for learnable_params_  
vector<float> params_lr_  
vector<bool> has_params_lr_  
/// the weight decay multipliers for learnable_params_  
vector<float> params_weight_decay_  
vector<bool> has_params_decay_  
/// The bytes of memory used by this net  
size_t memory_used_  
/// Whether to compute and display debug info for the net.  
bool debug_info_  
/// The root net that actually holds the shared layers in data parallelism  
const Net* const root_net_;
```

Solver会读取这些数值来控制对整个网络的更新, 比如, 某层的learning rate设0, 它在训练中就被冻结了, 参数不再改变。

Net的核心方法—当然也是向前向后方法

Net的前后方法

```
/**
 * The From and To variants of Forward and Backward operate on the
 * (topological) ordering by which the net is specified. For general DAG
 * networks, note that (1) computing from one layer to another might entail
 * extra computation on unrelated branches, and (2) computation starting in
 * the middle may be incorrect if all of the layers of a fan-in are not
 * included.
 */
const vector<Blob<Dtype>*>& Forward(Dtype* loss = NULL)
Dtype ForwardFromTo(int start, int end);
Dtype ForwardFrom(int start);
Dtype ForwardTo(int end);

/**
 * The network backward should take no input and output, since it solely
 * computes the gradient w.r.t the parameters, and the data has already been
 * provided during the forward pass.
 */
void Backward();
void BackwardFromTo(int start, int end);
void BackwardFrom(int start);
void BackwardTo(int end);
```

Net的“前后”方法和向前方法的实现

```
Dtype ForwardBackward() {
    Dtype loss;
    Forward(&loss);
    Backward();
    return loss;
}

template <typename Dtype>
const vector<Blob<Dtype>*>& Net<Dtype>::Forward(Dtype* loss) {
    if (loss != NULL) {
        *loss = ForwardFromTo(0, layers_.size() - 1);
    } else {
        ForwardFromTo(0, layers_.size() - 1);
    }
    return net_output_blobs_;
}

template <typename Dtype>
Dtype Net<Dtype>::ForwardFromTo(int start, int end) {
    CHECK_GE(start, 0);
    CHECK_LT(end, layers_.size());
    Dtype loss = 0;
    for (int i = start; i <= end; ++i) {
        // LOG(ERROR) << "Forwarding " << layer_names_[i];
        Dtype layer_loss = layers_[i]->Forward(bottom_vecs_[i], top_vecs_[i]);
        loss += layer_loss;
        if (debug_info_) { ForwardDebugInfo(i); }
    }
    return loss;
}
```

Net的向后方法

Backward 实际是由BackwardFromTo 实现的

```
template <typename Dtype>
void Net<Dtype>::BackwardFromTo(int start, int end) {
    CHECK_GE(end, 0);
    CHECK_LT(start, layers_.size());
    for (int i = start; i >= end; --i) {
        if (layer_need_backward_[i]) {
            layers_[i]->Backward(
                top_vecs_[i], bottom_need_backward_[i], bottom_vecs_[i]);
            if (debug_info_) { BackwardDebugInfo(i); }
        }
    }
}
```

Net提供的Update

```
template <typename Dtype>
void Net<Dtype>::Update() {
    for (int i = 0; i < learnable_params_.size(); ++i) {
        learnable_params_[i]->Update();
    }
}
```


最naive的方式就是：走完一个向前向后过程后，Blob里就有 $\frac{\partial l}{\partial \text{this blob}}$ 的数据，调用上面的Net的update方法更新参数。后面我们会看到有些Solver并不调用Net的Update，而是根据某些修正，去直接修改parameter的值。

另一点值得注意的是，最终的Forward的计算还是要依赖底层Layer的实现的，下面是具体的例子：Leaky ReLU。

Layer Forward/Backward的例子-Leaky ReLU定义

Leaky ReLU: 和普通ReLU的差别是小于0部分的斜率不再是0，而是一个可学习的参数 α ，反馈信号可以 α 比例“漏”过去。公式：

$$Y = X_+ - \alpha X_-$$



Layer Forward/Backward的例子—Leaky ReLU

向前实现: CPU版本

```
template <typename Dtype>
void ReLULayer<Dtype>::Forward_cpu(const vector<Blob<Dtype>*>& bottom,
    const vector<Blob<Dtype>*>& top) {
    const Dtype* bottom_data = bottom[0]->cpu_data();
    Dtype* top_data = top[0]->mutable_cpu_data();
    const int count = bottom[0]->count();
    Dtype negative_slope = this->layer_param_.relu_param().negative_slope();
    for (int i = 0; i < count; ++i) {
        // This is the line do the computation
        top_data[i] = std::max(bottom_data[i], Dtype(0))
            + negative_slope * std::min(bottom_data[i], Dtype(0));
    }
}
```

相对应的向前GPU Kernel

```
template <typename Dtype>
__global__ void ReLUForward(const int n, const Dtype* in, Dtype* out,
    Dtype negative_slope) {
    CUDA_KERNEL_LOOP(index, n) {
        out[index] = in[index] > 0 ? in[index] : in[index] * negative_slope;
    }
}
```

Layer Forward/Backward的例子—Leaky ReLU

向后实现: CPU版本

```
template <typename Dtype>
void ReLULayer<Dtype>::Backward_cpu(const vector<Blob<Dtype>*>& top,
    const vector<bool>& propagate_down,
    const vector<Blob<Dtype>*>& bottom) {
    if (propagate_down[0]) {
        const Dtype* bottom_data = bottom[0]->cpu_data();
        const Dtype* top_diff = top[0]->cpu_diff();
        Dtype* bottom_diff = bottom[0]->mutable_cpu_diff();
        const int count = bottom[0]->count();
        Dtype negative_slope = this->layer_param_.relu_param().negative_slope();
        for (int i = 0; i < count; ++i) {
            // the local gradient is 1 or negative_slope depend on X>0
            bottom_diff[i] = top_diff[i] * ((bottom_data[i] > 0) + negative_slope * (
                bottom_data[i] <= 0));
        } } }
```

相对应的向后GPU Kernel

```
template <typename Dtype>
__global__ void ReLUBackward(const int n, const Dtype* in_diff,
    const Dtype* in_data, Dtype* out_diff, Dtype negative_slope) {
    CUDA_KERNEL_LOOP(index, n) {
        out_diff[index] = in_diff[index] * ((in_data[index] > 0)
            + (in_data[index] <= 0) * negative_slope);
    } }
```

目录

1 显卡计算简介

2 Caffe框架解析

- Blob—数据对象
- Layer—变换层
- Net—整体网络
- Solver—求解器

3 TensorFlow—兵器谱之首位

- TensorFlow为什么叫TensorFlow?
- TensorFlow进行简单计算例子
- Perceptron的例子

求解器

求解器的用处

Net 定义了Hypotheses的结构，统计上相当于某种先验的知识。它的向前方法负责计算loss，向后方法计算反馈。而求解器负责如何从数据中取出一个个batch，通过计算loss和反馈来调整参数值，使得Net 收敛到一个我们满意的结果。

Solver类的核心域

```
SolverParameter param_;  
int iter_;  
int current_step_;  
shared_ptr<Net<Dtype> > net_;  
vector<shared_ptr<Net<Dtype> > > test_nets_;  
vector<Callback*> callbacks_;  
vector<Dtype> losses_;  
Dtype smoothed_loss_;
```

这里面控制求解器行为的最重要参数是param_，它的类SolverParameter 是由Protobuf定义的。我们下面详细解释。

SolverParameter的ProtoBuf定义

learning policy

```
// Proto filename for the train net, possibly combined with one or more test nets.
optional string net = 24;
optional float base_lr = 5; // The base learning rate
optional int32 max_iter = 7; // the maximum number of iterations

// The learning rate decay policy. The currently implemented learning rate policies
// are as follows:
// - fixed: always return base_lr.
// - step: return base_lr * gamma ^ (floor(iter / step))
// - exp: return base_lr * gamma ^ iter
// - inv: return base_lr * (1 + gamma * iter) ^ (- power)
// - multistep: similar to step but it allows non uniform steps defined by
//   stepvalue
// - poly: the effective learning rate follows a polynomial decay, to be
//   zero by the max_iter. return base_lr (1 - iter/max_iter) ^ (power)
// - sigmoid: the effective learning rate follows a sigmoid decay
//   return base_lr ( 1/(1 + exp(-gamma * (iter - stepsize))) )
//
// where base_lr, max_iter, gamma, step, stepvalue and power are defined
// in the solver parameter protocol buffer, and iter is the current iteration.
optional string lr_policy = 8;
optional float gamma = 9; // The parameter to compute the learning rate.
optional float power = 10; // The parameter to compute the learning rate.
optional int32 stepsize = 13; // the stepsize for learning rate policy "step"
repeated int32 stepvalue = 34; // the stepsize for learning rate policy "multistep"
```

SolverParameter的ProtoBuf定义

Regularization

```
optional float weight_decay = 12; // The weight decay.  
// regularization types supported: L1 and L2  
// controlled by weight_decay  
optional string regularization_type = 29 [default = "L2"];  
  
// Set clip_gradients to >= 0 to clip parameter gradients to that L2 norm,  
// whenever their actual L2 norm is larger.  
optional float clip_gradients = 35 [default = -1];
```

SolverParameter的ProtoBuf定义

Solver Type and Solver related parameters

```
// type of the solver: SGD, NESTEROV, ADAGRAD, RMSPROP, ADADELTA, ADAM
optional string type = 40 [default = "SGD"];

optional float momentum = 11; // The momentum value.

// numerical stability for RMSProp, AdaGrad and AdaDelta and Adam
optional float delta = 31 [default = 1e-8];

// parameters for the Adam solver
optional float momentum2 = 39 [default = 0.999];

// RMSProp decay value
// MeanSquare(t) = rms_decay*MeanSquare(t-1) + (1-rms_decay)*SquareGradient(t)
optional float rms_decay = 38;
```


求解器的核心方法是Step

因为Solve 方法的核心只有下面一句：

```
Step(param_.max_iter() - iter_);
```

Step里面的loop，每个loop是一个iteration。由于DataLayer也是Layer，它也有Forward方法（当然没有Backward）。所以每次iteration是由DataLayer负责去读取一个新的batch到data Blob中，再由实际计算的Layer向上传。

Step方法

```
template <typename Dtype> void Solver<Dtype>::Step(int iters) {
    const int start_iter = iter_;
    const int stop_iter = iter_ + iters;
    int average_loss = this->param_.average_loss();
    losses_.clear();
    smoothed_loss_ = 0;

    while (iter_ < stop_iter) {
        // zero-init the params
        net_->ClearParamDiffs();

        Dtype loss = 0;
        for (int i = 0; i < param_.iter_size(); ++i) {
            loss += net_->ForwardBackward();
        }
        loss /= param_.iter_size();
        // average the loss across iterations for smoothed reporting
        UpdateSmoothedLoss(loss, start_iter, average_loss);

        ApplyUpdate();
        ++iter_;
    }
}
```

ApplyUpdate 方法

Step 里面用向前向后来计算每个系数相关的更新量，放入Blob的diff中；然而具体如何使用 dW 去更新 W ，是依赖于具体方法实现的。

现有代码有下面的衍生类实现了相应的ApplyUpdate 方法：

- **SGDSolver** 带momentum的随机梯度下降
- **NesterovSolver** 用未来位置梯度更新的变种SGD
- **AdaGradSolver** 拟二阶的方法
- **RMSPropSolver** 相比AdaGrad，用指数平滑代替累加
- **AdaDeltaSolver** 相比AdaGrad，同时对梯度和参数改变做修正
- **AdamSolver** 同时使用一、二阶矩来光滑化修正。

我们下面一一解释：

SGDSolver –带动量的随机梯度下降

SGDSolver::ApplyUpdate 的核心部分就是

```
ClipGradients();  
for (int param_id = 0; param_id < this->net_->learnable_params().size();  
    ++param_id) {  
    Normalize(param_id);  
    Regularize(param_id);  
    ComputeUpdateValue(param_id, rate);  
}  
this->net_->Update();
```

前面计算好修正量，写回diff 部分，用Net 的update 方法更新之。

SGDSolver::ComputeUpdateValue 实际计算是由下面完成的

```
caffe_cpu_axpby(net_params[param_id]->count(), local_rate,  
               net_params[param_id]->cpu_diff(), momentum,  
               history_[param_id]->mutable_cpu_data());  
caffe_copy(net_params[param_id]->count(),  
          history_[param_id]->cpu_data(),  
          net_params[param_id]->mutable_cpu_diff());
```

对应的公式是：

$$\begin{aligned}\text{cache} &\leftarrow \text{local_rate} * dW + \text{momentum} * \text{cache} \\ dW &\leftarrow \text{cache}\end{aligned}$$

NesterovSolver—“一秒之后的未来”

用 $v^{(t)}$ 代表history，用 $\theta^{(t)}$ 代表迭代的当前值。用 r 代表learning rate而momentum的比例记作 μ 。上一页的公式就是下面

SGDSolver

$$\begin{aligned}v^{(t)} &\leftarrow r \cdot \nabla_{\theta} f(\theta^{(t-1)}) + \mu \cdot v^{(t-1)} \\ \theta^{(t)} &\leftarrow \theta^{(t-1)} - v^{(t)}\end{aligned}$$

NesterovSolver: 使用一步之后的未来梯度

$$\begin{aligned}v^{(t)} &\leftarrow r \cdot \nabla_{\theta} f(\theta^{(t-1)} - \mu v^{(t-1)}) + \mu \cdot v^{(t-1)} \\ \theta^{(t)} &\leftarrow \theta^{(t-1)} - v^{(t)}\end{aligned}$$

方便起见，把括号中用 $\phi^{(t-1)} = \theta^{(t-1)} - \mu v^{(t-1)}$ 代替，那么

$$\begin{aligned}v^{(t)} &\leftarrow r \cdot \nabla_{\theta} f(\phi^{(t-1)}) + \mu \cdot v^{(t-1)} \\ \phi^{(t)} &\leftarrow \phi^{(t-1)} + \mu v^{(t-1)} - (1 + \mu) v^{(t)}\end{aligned}$$

带二阶信息的方法—AdaGrad和RMSPropSolver

前面方法用到 ∇f ，如果引入 $(\nabla f)^2$ （点态），那么就有

AdaGradSolver

$$v^{(t)} \leftarrow v^{(t-1)} + (\nabla_{\theta} f)^2$$

$$\theta^{(t)} \leftarrow \theta^{(t-1)} - r \cdot \nabla_{\theta} f / (\sqrt{v^{(t)}} + \epsilon)$$

修正的意义相当于利用方差标准化。

RMSPropSolver

$$v^{(t)} \leftarrow \mu \cdot v^{(t-1)} + (1 - \mu) \cdot (\nabla_{\theta} f)^2$$

$$\theta^{(t)} \leftarrow \theta^{(t-1)} - r \cdot \nabla_{\theta} f / (\sqrt{v^{(t)}} + \epsilon)$$

用指数平滑代替了AdaGrad中的累加

带二阶信息的方法—AdaDeltaSolver

AdaDeltaSolver

引入 $E[g^2]_t$ 为 t 时刻的gradient g^2 估计, $E[\Delta x^2]_t$ 为 t 时刻参数改变量的估计, 迭代公式是:

$$E[g^2]_t \leftarrow \mu E[g^2]_{t-1} + (1 - \mu) g_t^2$$

$$\Delta x_t \leftarrow -\frac{\text{RMS}[\Delta x]_{t-1}}{\text{RMS}[g]_t} g_t$$

$$E[\Delta x^2]_t \leftarrow \mu E[\Delta x^2]_{t-1} + (1 - \mu) \Delta x_t^2$$

$$\text{where } \text{RMS}[y]_t = \sqrt{E[y^2]_t + \epsilon}$$

$$x_{t+1} \leftarrow x_t + r \cdot \Delta x_t$$

含义是用标准化了的 $g_t/\sigma(g)$ 去修正标准化的 $\Delta x_t/\sigma(\Delta x)$ 。

同时进行一二阶信息的平滑—AdamSolver

AdamSolver

$$m_t \leftarrow \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$v_t \leftarrow \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

$$x_t \leftarrow x_{t-1} - r \cdot \frac{m_t}{1 - \beta_1^t} \cdot \frac{1}{\sqrt{\frac{v_t}{1 - \beta_2^t}} + \epsilon}$$

实际使用中，基本上就先选**AdamSolver**就好了。如果它收敛不好，再试弱一点但**safer**一点的二阶或是一阶方法。

AdamSolver的修正量的解释

这里平滑化是很自然的想法，唯一要小心的是修正。以 m_t 为例，展开之

$$m_t = \beta_1^{t-1}(1 - \beta_1)g_1 + \beta_1^{t-2}(1 - \beta_1)g_2 + \cdots + (1 - \beta_1)g_t$$

两边取Expectation，有


$$\begin{aligned} E[m_t] &= (\beta_1^{t-1}(1 - \beta_1) + \beta_1^{t-2}(1 - \beta_1) + \cdots + (1 - \beta_1))E[g_t] \\ &= (1 - \beta_1) \frac{1 - \beta_1^t}{1 - \beta_1} E[g_t] \\ &= (1 - \beta_1^t) E[g_t] \end{aligned}$$

即

$$E[g_t] = \frac{E[m_t]}{1 - \beta_1^t}$$

完整的例子

LeNet



而传给Solver的参数如下:

```
net: "examples/mnist/lenet_train_test.prototxt"
test_iter: 100
test_interval: 500
base_lr: 0.01
momentum: 0.9
weight_decay: 0.0005
lr_policy: "inv"
gamma: 0.0001
power: 0.75
# Display every 100 iterations
display: 100
max_iter: 10000
snapshot: 5000
snapshot_prefix: "examples/mnist/lenet"
solver_mode: GPU
```

训练命令行:

```
./build/tools/caffe train --solver=examples/mnist/lenet_solver.  
prototxt
```

Caffe的问题

- **开发难度高** 新的Layer、Loss一般需要写C++和CUDA代码才能在Caffe中使用；虽然有Python-Layer，可以使用Python来写Layer代码，但它只实现了CPU部分，会拖慢整个的计算流水线。
- **Backward开发难** 需要手动推导local gradient的公式，才能实现代码
- **搭建Net不灵活** 复杂的Net结构需要写非常复杂的protobuf描述文件。
- **修改Layer不灵活** 微调某些层的公式只有修改Layer的实现类的代码，重新编译。

如何写出优雅的代码？




答案：

用更加合适的工具


下面的图可能带来一些启发！

流行的深度框架热度对比


Aggregate popularity (30*contrib + 10*issues + 5*forks)*1e-3




Top libraries by Github issues opened




Top libraries by Github stars



Top libraries by Github contributors



Top libraries by Github forks



目录

- 1 显卡计算简介
- 2 Caffe框架解析
 - Blob—数据对象
 - Layer—变换层
 - Net—整体网络
 - Solver—求解器
- 3 TensorFlow—兵器谱之首位
 - **TensorFlow**为什么叫TensorFlow?
 - TensorFlow进行简单计算例子
 - Perceptron的例子

TensorFlow的本质

从上一讲，DL里的构造Net的问题本质上只有两个：

- 简单表达式计算的向量（张量）化：比如 $\exp(a) + b$ ，这里 a, b 都是向量
- 向量化的复合映射的求导

解决了这两个问题，就解决了BP算法中最核心的张量向前向后流动的问题，也就解决了编写Layer困难的问题。而大G站的TensorFlow正是解决这类编程问题的很好的工具。

顾名思义

Tensor(张量)**Flow**(流动)—正是解决张量计算和自动求导的工具

有了上面两个特性，我们甚至可以用TensorFlow解很多和DeepLearning无关的问题，比如传统的ML甚至一些优化问题。

要优雅

TensorFlow是如何优雅地解决这两个问题的？

解决方法

Computational Graph

TensorFlow可以看作是描述Computational Graph的DSL

下面看具体的例子，TensorFlow的每一次实际执行都被组织成一个个session，我们先准备一个session。

准备

```
from tensorflow import *  
import numpy as np  
sess = Session()
```

目录

- 1 显卡计算简介
- 2 Caffe框架解析
 - Blob—数据对象
 - Layer—变换层
 - Net—整体网络
 - Solver—求解器
- 3 TensorFlow—兵器谱之首位
 - TensorFlow为什么叫TensorFlow?
 - TensorFlow进行简单计算例子
 - Perceptron的例子

Affine变换：矩阵乘法

一个简单的矩阵乘法

$$X = \begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{pmatrix} \quad W = \begin{pmatrix} 1 \\ 2 \\ 3 \end{pmatrix} \quad Y = X \cdot W$$

定义张量

```
X = constant([[1,2,3],[4,5,6],[7,8,9]],name='X')  
W = constant([1,2,3], shape=[3,1], name='W')
```

矩阵乘法

```
Y = matmul(X, W, name = 'Y')
```

计算结果

```
Y0 = sess.run(Y)  
> [[14]  
> [32]  
> [50]]
```

Affine变换：矩阵乘法的梯度

梯度张量 $\frac{\partial Y}{\partial X}$

$$\left[\frac{\partial Y_i}{\partial X_{k,j}} \right]_{i,k,j} = \delta_{ik} \cdot W_j$$

梯度张量 $\frac{\partial Y}{\partial W}$

$$\left[\frac{\partial Y_i}{\partial W_j} \right]_{i,j} = X_{ij} = X$$

TensorFlow里gradient计算给出的是 $\sum_i \frac{\partial Y_i}{\partial \theta_j}$ 。

对上面的例子，结果是 $\mathbf{1} \cdot W^T$ 和 $X^T \cdot \mathbf{1}$ 。

Affine变换：矩阵乘法的梯度

gradients 能通过链式法则进行符号计算，求出相应的梯度。

求梯度

```
dX = gradients(Y, X, name='dX')
dW = gradients(Y, W, name='dW')
sess.run(dX[0])
>array([[1, 2, 3],
>        [1, 2, 3],
>        [1, 2, 3]], dtype=int32)

sess.run(dW[0])
>array([[12],
>        [15],
>        [18]], dtype=int32)
```

检验结果

```
II = ones((3,1),int32)
sess.run(reduce_all(equal(dX[0] , matmul(II, transpose(W)))))
>True
sess.run(reduce_all(equal(dW[0] , matmul(transpose(X) , II))))
>True
```

ReLU – maximum

maximum 做elementwise的比较，返回大的；而且它正确处理了反向求gradient的计算。

local gradients

```
Z = constant([-1,2,0], shape=[3,1], name='Z')
sess.run(gradients(maximum(Z,0), Z))
>[array([[0],
>         [1],
>         [1]], dtype=int32)]
```

注意

0点不可导，TensorFlow的处理，就是令0点的导数为一。

上方有回传的gradient时，ReLU相当于一个开关

```
sess.run(gradients(maximum(Z,0), Z, grad_ys=constant([100,99,98], shape=(3,1))))
>[array([[ 0],
>         [99],
>         [98]], dtype=int32)]
```

上层传来的gradient是(100,99,98)但只有2,3位置通过了。

知足而常乐

毕竟这里的符号计算只是规则引擎，能处理链式法则，但不能真的求解数学问题，比如：

隐函数/参数表示的求导

$$x^2 + y^2 = 1 \quad \text{or} \quad \begin{cases} x = \cos(t) \\ y = \sin(t) \end{cases}$$

那么，

$$\frac{dy}{dx} = \frac{-\sin(t) dt}{\cos(t) dt} = -\frac{x}{y}$$

实验：参数方程求gradient

```
t = constant(1.0)
x = cos(t)
y = sin(t)
gradients(y, x)
> [None]
```

显化：对比正确值

```
sess.run(gradients(sin(acos(x)), x))
> [-0.64209253]


sess.run(-x/y)
> -0.64209259
```

目录

- 1 显卡计算简介
- 2 Caffe框架解析
 - Blob—数据对象
 - Layer—变换层
 - Net—整体网络
 - Solver—求解器
- 3 TensorFlow—兵器谱之首位
 - TensorFlow为什么叫TensorFlow?
 - TensorFlow进行简单计算例子
 - Perceptron的例子

Perceptron例子：问题

问题:平面上找直线分离正例和负例



定义数据和参数

```
# training data
X = constant([[5,1],[6,1],[1,3],[2,4],[3,5]], name='data', dtype=float32)
Y = constant([1, 1, -1, -1, -1], shape=[5,1], name='label', dtype=float32)
# parameters
W = Variable(zeros((2,1),dtype=float32))
b = Variable(1, dtype=float32)
```

Perceptron例子: Loss函数

使用maximal margin的Loss

Score是

$$S = X \cdot W + b$$

我们希望在 $Y = 1$ 的score越大越好, $Y = -1$ 的score越小越好, 所以Loss是

$$l = (1 - S_{Y>0})_+ + (S_{Y<0} + 1)_+$$

最终分类器的输出是 $\text{sign}(S)$, 所以训练集上Accuracy是

$\{i | \text{sign}(S_i) = \text{sign}(Y_i)\}$ 的个数/样本数量

Score, Loss和正确率的计算公式

```
# compute score, loss and accuracy
S = matmul(X, W) + b
loss = reduce_sum(maximum(1-S[0:1,0],0)) + reduce_sum(maximum(S[2:4,0]+1,0))
accuracy = reduce_mean(to_float(equal(sign(S), sign(Y))))
```

Perceptron例子：求解

回忆：最简单的参数训练方式

$$W \leftarrow W - r \cdot dW$$

这里 r 是learning rate。

定义BP的update rule

```
#learning rate
lr = 0.1
#updating rule
dW, db = gradients(loss, [W,b])
update_param = group( W.assign_sub( lr * dW), b.assign_sub( lr * db ) )
```

梯度下降求解

```
it = 0
while True:
    it = it + 1
    l, acc = sess.run([loss, accuracy])
    print "Iter : %d Loss = %.2f Accuracy=%.2f%%\n" % (it, l, acc*100)
    if (l<0.05):
        break
    sess.run(update_param)
```

Perceptron结果

结果

```
>Iter : 1 Loss = 4.00 Accuracy=40.00%
>
>Iter : 2 Loss = 0.30 Accuracy=100.00%
>
>Iter : 3 Loss = 0.60 Accuracy=100.00%
>
>Iter : 4 Loss = 0.10 Accuracy=100.00%
>
>Iter : 5 Loss = 0.00 Accuracy=100.00%

sess.run(S)

>array([[ 2.19999981],
>       [ 2.69999981],
>       [-2.00000024],
>       [-2.60000014],
>       [-3.20000005]], dtype=float32)

sess.run([W,b])


>[array([[ 0.5          ],
>       [-1.10000002]], dtype=float32), 0.79999995]
```

Perceptron结果：可视化

分离直线的方程

从上面的 $W = (0.5, -1.1)^T$, $b = 0.8$, 我们知道方程是

$$0.5x_1 - 1.1x_2 + 0.8 = 0$$



可以看到我们选取的**Loss**函数确实使得直线落在比较好的位置。

Perceptron: 直接使用优化器

优化器(Optimizer)类可以产生参数的updater, 可以取代手写的update_param 。

使用GradientDescentOptimizer取代update_param

```
opt = train.GradientDescentOptimizer(lr)
update_using_opt = opt.minimize(loss)
it = 0
while True:
    it = it + 1
    l, acc = sess.run([loss, accuracy])
    print "Iter : %d Loss = %.2f Accuracy=%.2f%%\n" % (it, l, acc*100)
    if (l<0.05):
        break
    sess.run(update_using_opt)

>Iter : 1 Loss = 4.00 Accuracy=40.00%
>
>Iter : 2 Loss = 0.30 Accuracy=100.00%
>
>Iter : 3 Loss = 0.60 Accuracy=100.00%
>
>Iter : 4 Loss = 0.10 Accuracy=100.00%
>
>Iter : 5 Loss = 0.00 Accuracy=100.00%
```

结果也是一样的。

Perceptron: 使用cross entropy loss

使用softmax_cross_entropy_with_logits

```
X = constant([[5,1],[6,1],[1,3],[2,4],[3,5]], name='data', dtype=float32)
# label changed to one_hot format
label_one_hot = one_hot([0,0,1,1,1], 2)
W = Variable(zeros((2,1),dtype=float32))
b = Variable(1, dtype=float32)

S = matmul(X, W) + b
# score need to be [S,1] format
loss = reduce_mean(nn.softmax_cross_entropy_with_logits(pad(S, [[0,0],[0,1]]),
    label_one_hot, name='cross_entropy'))
# use sign(label_one_hot - 0.5) generate +/-1
accuracy = reduce_mean(to_float(equal(sign(S) , sign(matmul(label_one_hot, constant
    ([1.0,-1.0],shape=(2,1)))))))

#following is same as previous
#learning rate
lr = 0.1
#updating rule
opt = train.GradientDescentOptimizer(lr)
update_using_opt = opt.minimize(loss)
it = 0
while True:
    it = it + 1
    l, acc = sess.run([loss, accuracy])
    print "Iter : %d Loss = %.2f Accuracy=%.2f%%\n" % (it, l, acc*100)
    if (l<0.05):
        break
    sess.run(update_using_opt)
```

Perceptron: 使用cross entropy loss的结果

结果

```
>Iter : 1 Loss = 0.91 Accuracy=40.00%
>
>Iter : 2 Loss = 0.67 Accuracy=40.00%
>
>Iter : 3 Loss = 0.54 Accuracy=80.00%
>
>Iter : 4 Loss = 0.45 Accuracy=100.00%
>
> . . .
>
>Iter : 47 Loss = 0.05 Accuracy=100.00%
>
>Iter : 48 Loss = 0.05 Accuracy=100.00%

sess.run(S)


>array([[ 2.71125937],
>       [ 3.34846902],
>       [-2.4624207 ],
>       [-3.13763237],
>       [-3.81284356]], dtype=float32)

sess.run([W,b])

>[array([[ 0.63720953],
>       [-1.31242108]], dtype=float32), 0.83763283]
>
```

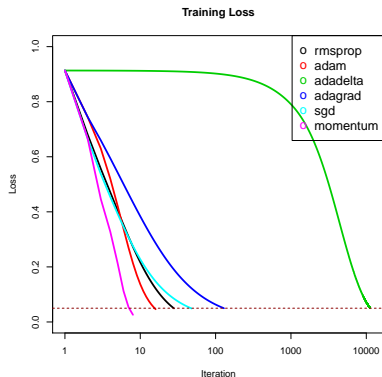

Perceptron: 使用cross entropy loss的结果

$$0.6372x_1 - 1.3124x_2 + 0.8376 = 0$$



虽然迭代次数多了（因为loss的尺度不同，不能直接比较），最终效果和SVM loss差不多

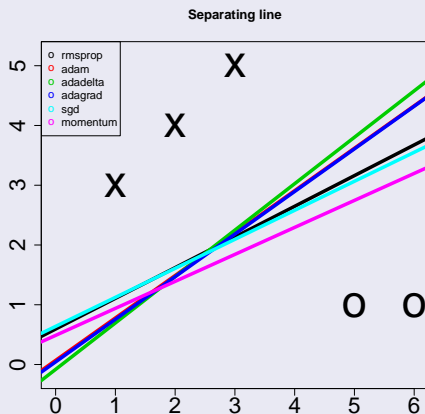
Perceptron: 比较不同优化器的结果



	k	b	迭代次数
rmsprop	0.51	0.59	28.00
adam	0.71	0.07	16.00
adadelata	0.78	-0.08	11241.00
adagrad	0.71	0.04	129.00
sgd	0.49	0.64	48.00
momentum	0.45	0.49	8.00

Perceptron比较不同优化器的结果

分离超平面



课后的欢乐小剧场



“他强由他强，清风拂山岗；他横由他横，明月照大江。他自狠来他自恶，我自一口真气足。。。 ”

“姐姐姐姐，人家是感慨一下理解了基础原理多么重要啊”



“什么鬼，雷姆雷姆，你听课听地画风都变啦”

下周预告



“回顾一下，今天讨论了

- GPU计算的概念和困难
- Caffe框架的结构
- TensorFlow的开发理念

好像有不少收获呢！”

“还是姐姐的魔法最棒了！
下次我们将开始这场旅程
最后的探险，哦。。。 ”



“DeepLearning看起来
也不是那么神秘嘛！比
起魔法差多了。。。 ”